

```
In [1]: import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from sklearn.decomposition import PCA
import torch.utils.data
import pickle
from tqdm import tqdm
import time
import gc
import collections
```

```
In [2]: '''
set parameters
'''

base_path = '.'
train_path = base_path + '/train.npy'
test_path = base_path + '/test.npy'
train_labels_path = base_path + '/train_labels.npy'
dev_labels_path = base_path + '/dev_labels.npy'
dev_path = base_path + '/dev.npy'

pca_available = True
padding_method = 'self'
# padding_method = 'zero'

device = torch.device('cuda:0')
n_labels = 138
n_features = 40
n_epoch = 10
context_num = 18
```

```
In [3]: def load_train_data():
    t0 = time.time()
    print("Start loading training data...")
    train = np.load(train_path, allow_pickle=True)
    train_labels = np.load(train_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading training data in {0} minutes...".format(elapsed_time/60))

    return train, train_labels
```

```
In [4]: def load_validation_data():
    t0 = time.time()
    print("Start loading validation data...")
    val = np.load(dev_path, allow_pickle=True)
    val_labels = np.load(dev_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading validation data in {0} minutes...".format(elapsed_time/60))

    return val, val_labels
```

```
In [5]: def load_test_data():
    t0 = time.time()
    print("Start loading test data...")
    test = np.load(test_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading test data in {0} minutes...".format(elapsed_time/60))
    return test
```

```
In [6]: def load_and_process_data(features, labels, pca, context_num):
    """
    use the first&last frame of one utterance to pad the empty frame
    """
    t0 = time.time()
    padding_features = np.concatenate([np.concatenate(( \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=np.float32), \
                                                pca.transform(features[i]), \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=np.float32)) \
                                                for i in range(len(features)))]
    padding_features = torch.Tensor(padding_features)
    del features

    """
    corresponding label for padding frames
    """
    false_labels = np.array([-1]*context_num)
    padding_labels = np.concatenate([np.concatenate(( \
                                                false_labels, \
                                                labels[i], \
                                                false_labels)) \
                                                for i in range(len(labels)))]
    padding_labels = torch.Tensor(padding_labels)
    del labels

    gc.collect()

    return padding_features, padding_labels
```

```

In [7]: ass ContextDataset(Dataset):

    def __init__(self, context_num, features, targets):

        self.context_num = context_num
        self.features = features
        self.targets = targets

    def __len__(self):

        return len(self.targets)

    def __getitem__(self, index):

        if index-self.context_num >= 0 and index+self.context_num+1 <= len(self.targets):
            '''
            no need for padding
            '''
            X = self.features[index-self.context_num: index+self.context_num+1].res
            Y = self.targets[index].long()
        elif index-self.context_num < 0:
            '''
            padding for pre frames, actually doesnt matter since we drop this 'false
            '''
            X = torch.cat((torch.zeros(self.context_num-index, self.features.shape[1]),
                           self.features[index-self.context_num:]).res
                           self.features[index-self.context_num:]).res
            Y = self.targets[index].long()
        else:
            '''
            padding for post frames, same as before
            '''
            X = torch.cat((self.features[index-self.context_num:], torch.zeros(self.context_num-index, self.features.shape[1])).res
                           self.features[index-self.context_num:]).res
            Y = self.targets[index].long()

        return index, X, Y

```

```

In [8]: class SpeechDataset(Dataset):

    def __init__(self, speechdataset):

        self.features = [speechdataset[i][1] for i in range(len(speechdataset))]
        self.targets = [speechdataset[i][2] for i in range(len(speechdataset))]

    def __len__(self):

        return len(self.targets)

    def __getitem__(self, index):

        return index, self.features[index], self.targets[index]

```

```

In [9]: #1000,2048,1024,512,256+2,138
class SpeechNet(nn.Module):

    def __init__(self, context_num):

        super(SpeechNet, self).__init__()
        self.relu1 = nn.ReLU()
        self.relu2 = nn.ReLU()
        self.relu3 = nn.ReLU()
        self.relu4 = nn.ReLU()
        self.relu5 = nn.ReLU()
        self.relu6 = nn.ReLU()
        self.relu7 = nn.ReLU()
        self.relu8 = nn.ReLU()

        self.linear1 = nn.Linear((2*context_num+1)*pca.n_components, 2048)
        self.linear2 = nn.Linear(2048, 1024)
        self.linear3 = nn.Linear(1024, 810)
        self.linear4 = nn.Linear(810, 720)
        self.linear5 = nn.Linear(720, 512)
        self.linear6 = nn.Linear(512, 428)
        self.linear7 = nn.Linear(428, 300)
        self.linear8 = nn.Linear(300, 256)
        self.out = nn.Linear(256+2, 138)

        self.batchnorm1 = nn.BatchNorm1d(2048)
        self.batchnorm2 = nn.BatchNorm1d(1024)
        self.batchnorm3 = nn.BatchNorm1d(810)
        self.batchnorm4 = nn.BatchNorm1d(720)
        self.batchnorm5 = nn.BatchNorm1d(512)
        self.batchnorm6 = nn.BatchNorm1d(428)
        self.batchnorm7 = nn.BatchNorm1d(300)
        self.batchnorm8 = nn.BatchNorm1d(256+2)

        self.dropout1 = nn.Dropout(0.1)
        self.dropout2 = nn.Dropout(0.05)
        self.dropout3 = nn.Dropout(0.1)
        self.dropout4 = nn.Dropout(0.05)
        self.dropout5 = nn.Dropout(0.05)
        self.dropout6 = nn.Dropout(0.05)
        self.dropout7 = nn.Dropout(0.05)

    def forward(self, x):

        x = self.linear1(x)
        x = self.batchnorm1(x)
        x = self.relu1(x)
        #2048
        x = self.dropout1(x)
        x = self.linear2(x)
        x = self.batchnorm2(x)
        x = self.relu2(x)
        #1024
        x = self.dropout2(x)
        x = self.linear3(x)

```

```

x = self.batchnorm3(x)
x = self.relu3(x)
#810
x = self.dropout3(x)
x = self.linear4(x)
x = self.batchnorm4(x)
x = self.relu4(x)

#512
x = self.dropout4(x)
x = self.linear5(x)
x = self.batchnorm5(x)
x = self.relu5(x)

#512
x = self.dropout5(x)
x = self.linear6(x)
x = self.batchnorm6(x)
x = self.relu6(x)
#428
x = self.dropout6(x)
x = self.linear7(x)
x = self.batchnorm7(x)
x = self.relu7(x)
#300
x = self.dropout7(x)
x = self.linear8(x)
#300

avg_pool1 = torch.mean(x, 1, keepdims = True)
max_pool1,_ = torch.max(x, 1, keepdims = True)

conc = torch.cat((x, avg_pool1, max_pool1), 1)
conc = self.batchnorm8(conc)
output = self.out(conc)

return output

```

```

In [10]: def generate_dataset(context_num, features, labels):
    t0 = time.time()
    print("It may takes 20 minutes to generate train dataset...")
    context_dataset = ContextDataset(context_num, features, labels)
    dataset = SpeechDataset(context_dataset)
    t1 = time.time()
    print("Dataset generated. Elapsed time: {0}".format((t1-t0)/60))
    return dataset

```

```

In [11]: def weights_init(m):
    if isinstance(m, nn.Conv2d):
        xavier(m.weight.data)
        xavier(m.bias.data)

```

```
In [12]: def scale_cos(x):  
    start = 5e-3  
    end = 1e-5  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [13]: def second_scale_cos(x):  
    start = 1e-4  
    end = 1e-8  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [14]: class ParamScheduler:  
  
    def __init__(self, optimizer, scale_fn, total_steps):  
  
        self.optimizer = optimizer  
        self.scale_fn = scale_fn  
        self.total_steps = total_steps  
        self.current_iteration = 0  
  
    def batch_step(self):  
        for param_group in self.optimizer.param_groups:  
            param_group['lr'] = self.scale_fn(self.current_iteration/self.total_steps)  
  
        self.current_iteration += 1
```

```

In [15]: def train_model(train_dataloader, val_dataloader, n_epochs = 10):

    model = SpeechNet(context_num).to(device)
    model.apply(weights_init)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr = 0.001)

    '''
    set scheduler for decaying learning rate
    '''

    parameter_scheduler = ParamScheduler(optimizer, scale_cos, n_epoch*len(train_dataloader))
    candidate_model = 1
    print('Start training...')
    for i in range(n_epochs):

        t0 = time.time()

        avg_loss_1000_batch = 0
        val_correct = 0
        val_predicted = 0
        model.train()

        for index, (idx, features, labels) in enumerate(train_dataloader):

            optimizer.zero_grad()
            mask = [i for i in range(len(labels)) if labels[i] != torch.Tensor([0])]
            features = features[mask].cuda()
            labels = labels[mask].cuda()
            '''
            forward and backward
            '''

            output = model(features)
            loss = criterion(output, labels.long())
            avg_loss_1000_batch += loss.item()
            loss.backward()

            parameter_scheduler.batch_step()
            optimizer.step()

            if index % 2000 == 0 and index != 0:

                predictions = torch.max(output.data, 1)[1]
                predicted = len(features)
                correct = int(sum(predictions == labels.to(device)).cpu())
                print("Epoch: {0}/{1} Train batch:{2}/{3}  acc: {4}  loss: {5}"
                      .format(i, len(train_dataloader), len(features), len(labels), correct, loss.item()))

                avg_loss_1000_batch = 0

        for index, (idx, val_features, val_labels) in enumerate(val_dataloader):

            mask = [i for i in range(len(val_labels)) if val_labels[i] != torch.Tensor([0])]
            val_features = val_features[mask].to(device)
            val_labels = val_labels[mask].to(device)

```

```
model.eval()
outputs = model(val_features)
predictions = torch.max(outputs.data, 1)[1]
val_predicted += len(val_features)
val_correct += sum(predictions == val_labels.to(device))

epoch_acc = int(val_correct.cpu())/val_predicted
if epoch_acc >= 0.68:
    pickle.dump(model, open("candidate_model_{0}.pkl".format(candidate_model), "a"))
    print("Save one candidate model.")
    candidate_model += 1

t1 = time.time()
print("Validation Accuracy: {0}. Cost time: {1} minutes".format(int(val_correct.cpu()), t1 - t0))
print("=====")

return model
```



```

In [16]: if __name__ == '__main__':

    train_features, train_labels = load_train_data()
    val_features, val_labels = load_validation_data()
    test_features = load_test_data()

    if pca_available:
        '''
        load local pca file
        '''
        pca = pickle.load(open('pca.pkl', 'rb'))
    else:
        '''
        10 features will be enough
        '''
        pca = PCA(10).fit(np.concatenate(train_features))
        pickle.dump(pca, open('pca_{0}_features.pkl'.format(pca.n_components), 'w'))

    train_features, train_labels = load_and_process_data(train_features, train_labels)
    val_features, val_labels = load_and_process_data(val_features, val_labels, pca)

    train_context_dataset = ContextDataset(context_num, train_features, train_labels)
    val_context_dataset = ContextDataset(context_num, val_features, val_labels)

    '''
    It may takes more than 20 minutes to process since it loops over all 15 million samples
    But it could speed up later dataloader process
    And I could save this data by pickle
    The drawback is it's not flexible to feature engineering, like change context number
    '''
    '''
    train_dataset = SpeechDataset(train_context_dataset)
    val_dataset = SpeechDataset(val_context_dataset)
    '''
    '''
    train_mask = (train_dataset.targets.numpy() != -1)*1
    train_sampler = WeightedRandomSampler(weights=train_mask, num_samples=int(train_mask.sum()))
    val_mask = (val_dataset.targets.numpy() != -1)*1
    val_sampler = WeightedRandomSampler(weights=val_mask, num_samples=int(val_mask.sum()))
    '''
    '''
    train_dataloader = DataLoader(train_dataset,
                                  shuffle = True,
                                  batch_size = 512,
                                  num_workers = 0,
                                  pin_memory = True)

    val_dataloader = DataLoader(val_dataset,
                                shuffle = True,
                                batch_size = 512,
                                num_workers = 0,
                                pin_memory = True)

    '''
    '''
    Start train model.
    15 epochs by default.

```

Cost about 1 hour.

'''

```
train_dataloader = DataLoader(train_context_dataset,
                              shuffle = True,
                              batch_size = 512,
                              num_workers = 0,
                              pin_memory = True)

val_dataloader = DataLoader(val_context_dataset,
                             shuffle = True,
                             batch_size = 512,
                             num_workers = 0,
                             pin_memory = True)

model = train_model(train_dataloader, val_dataloader)

pickle.dump(model, open("submission_model_68.pkl", "wb"))
#make_submission(model, test_features)
```

Start loading training data...

Done loading training data in 0.7987695336341858 minutes...

Start loading validation data...

Done loading validation data in 0.0456453800201416 minutes...

Start loading test data...

Done loading test data in 0.014128851890563964 minutes...

Start training...

Epoch: 1/10 Train batch:2000/31779 acc: 0.45121951219512196 loss: 9.797406  
8808835

Epoch: 1/10 Train batch:4000/31779 acc: 0.4979253112033195 loss: 8.1310059  
98002365

Epoch: 1/10 Train batch:6000/31779 acc: 0.5040816326530613 loss: 7.6088751  
33330002

Epoch: 1/10 Train batch:8000/31779 acc: 0.5708418891170431 loss: 7.2945281  
34632856

Epoch: 1/10 Train batch:10000/31779 acc: 0.5343035343035343 loss: 7.054817  
696334794

Epoch: 1/10 Train batch:12000/31779 acc: 0.5429769392033543 loss: 6.877835  
999010131

Epoch: 1/10 Train batch:14000/31779 acc: 0.5979381443298969 loss: 6.748973  
261332139

Epoch: 1/10 Train batch:16000/31779 acc: 0.5637860082304527 loss: 6.626700  
401538983

Epoch: 1/10 Train batch:18000/31779 acc: 0.56875 loss: 6.515335530042648

Epoch: 1/10 Train batch:20000/31779 acc: 0.5379876796714579 loss: 6.443903  
14957127

Epoch: 1/10 Train batch:22000/31779 acc: 0.5938775510204082 loss: 6.355136  
304395273

Epoch: 1/10 Train batch:24000/31779 acc: 0.5967078189300411 loss: 6.303867  
116337642

Epoch: 1/10 Train batch:26000/31779 acc: 0.5983263598326359 loss: 6.232847  
114559263

Epoch: 1/10 Train batch:28000/31779 acc: 0.5479166666666667 loss: 6.187457  
967782393

Epoch: 1/10 Train batch:30000/31779 acc: 0.5860655737704918 loss: 6.145368  
884084746

Validation Accuracy: 0.6099275563894199. Cost time: 24.387852080663045 minute  
s

```
=====
Epoch: 2/10 Train batch:2000/31779 acc: 0.6356107660455487 loss: 6.0118769
84499395
Epoch: 2/10 Train batch:4000/31779 acc: 0.5904365904365905 loss: 5.9843146
83126286
Epoch: 2/10 Train batch:6000/31779 acc: 0.6290983606557377 loss: 5.9456812
27371097
Epoch: 2/10 Train batch:8000/31779 acc: 0.5785123966942148 loss: 5.9185198
9692077
Epoch: 2/10 Train batch:10000/31779 acc: 0.6042105263157894 loss: 5.886191
22421369
Epoch: 2/10 Train batch:12000/31779 acc: 0.6125 loss: 5.852636412251741
Epoch: 2/10 Train batch:14000/31779 acc: 0.5871369294605809 loss: 5.832868
226803839
Epoch: 2/10 Train batch:16000/31779 acc: 0.5857142857142857 loss: 5.808117
415057495
Epoch: 2/10 Train batch:18000/31779 acc: 0.5995893223819302 loss: 5.789125
464623794
Epoch: 2/10 Train batch:20000/31779 acc: 0.59375 loss: 5.754060934996232
Epoch: 2/10 Train batch:22000/31779 acc: 0.6470588235294118 loss: 5.732057
94300884
Epoch: 2/10 Train batch:24000/31779 acc: 0.6504065040650406 loss: 5.722947
325091809
Epoch: 2/10 Train batch:26000/31779 acc: 0.6164948453608248 loss: 5.697432
388318703
Epoch: 2/10 Train batch:28000/31779 acc: 0.6244813278008299 loss: 5.680661
244317889
Epoch: 2/10 Train batch:30000/31779 acc: 0.6149068322981367 loss: 5.653364
775236696
Validation Accuracy: 0.6388902041323635. Cost time: 24.138027866681416 minute
s
=====
Epoch: 3/10 Train batch:2000/31779 acc: 0.6597938144329897 loss: 5.5778227
88618505
Epoch: 3/10 Train batch:4000/31779 acc: 0.6701244813278008 loss: 5.5390974
03176129
Epoch: 3/10 Train batch:6000/31779 acc: 0.6453608247422681 loss: 5.5392883
79950449
Epoch: 3/10 Train batch:8000/31779 acc: 0.6278118609406953 loss: 5.5180723
93490002
Epoch: 3/10 Train batch:10000/31779 acc: 0.6440329218106996 loss: 5.506670
6200595945
Epoch: 3/10 Train batch:12000/31779 acc: 0.6053169734151329 loss: 5.500768
842641264
Epoch: 3/10 Train batch:14000/31779 acc: 0.6529774127310062 loss: 5.487650
30503273
Epoch: 3/10 Train batch:16000/31779 acc: 0.6459627329192547 loss: 5.484318
618895486
Epoch: 3/10 Train batch:18000/31779 acc: 0.6523517382413088 loss: 5.469916
230067611
Epoch: 3/10 Train batch:20000/31779 acc: 0.6515463917525773 loss: 5.443030
263995752
Epoch: 3/10 Train batch:22000/31779 acc: 0.6040816326530613 loss: 5.427898
246562108
Epoch: 3/10 Train batch:24000/31779 acc: 0.6561181434599156 loss: 5.423127
093818039
Epoch: 3/10 Train batch:26000/31779 acc: 0.639344262295082 loss: 5.4192846
```

20322287

Epoch: 3/10 Train batch:28000/31779 acc: 0.6378600823045267 loss: 5.411432  
175664231

Epoch: 3/10 Train batch:30000/31779 acc: 0.663135593220339 loss: 5.3958395  
22359893

Validation Accuracy: 0.6519155534774709. Cost time: 24.139771298567453 minute  
s

=====

Epoch: 4/10 Train batch:2000/31779 acc: 0.6257668711656442 loss: 5.3162106  
93912581

Epoch: 4/10 Train batch:4000/31779 acc: 0.6548856548856549 loss: 5.2943583  
80604535

Epoch: 4/10 Train batch:6000/31779 acc: 0.5907172995780591 loss: 5.2959450  
49962029

Epoch: 4/10 Train batch:8000/31779 acc: 0.6321353065539113 loss: 5.2873974  
24690425

Epoch: 4/10 Train batch:10000/31779 acc: 0.5917525773195876 loss: 5.273880  
596272647

Epoch: 4/10 Train batch:12000/31779 acc: 0.6605691056910569 loss: 5.276953  
265769407

Epoch: 4/10 Train batch:14000/31779 acc: 0.6401673640167364 loss: 5.265228  
853095323

Epoch: 4/10 Train batch:16000/31779 acc: 0.6384297520661157 loss: 5.254212  
558735162

Epoch: 4/10 Train batch:18000/31779 acc: 0.6346555323590815 loss: 5.241534  
168599173

Epoch: 4/10 Train batch:20000/31779 acc: 0.675564681724846 loss: 5.2273399  
3944712

Epoch: 4/10 Train batch:22000/31779 acc: 0.6567796610169492 loss: 5.215027  
284575626

Epoch: 4/10 Train batch:24000/31779 acc: 0.6880165289256198 loss: 5.212054  
678238928

Epoch: 4/10 Train batch:26000/31779 acc: 0.656964656964657 loss: 5.2214122  
3680228

Epoch: 4/10 Train batch:28000/31779 acc: 0.6869918699186992 loss: 5.199652  
779381722

Epoch: 4/10 Train batch:30000/31779 acc: 0.6632860040567952 loss: 5.195865  
701651201

Validation Accuracy: 0.6607150255387402. Cost time: 24.175887401898702 minute  
s

=====

Epoch: 5/10 Train batch:2000/31779 acc: 0.6170212765957447 loss: 5.0965385  
02257317

Epoch: 5/10 Train batch:4000/31779 acc: 0.654320987654321 loss: 5.11735978  
4897417

Epoch: 5/10 Train batch:6000/31779 acc: 0.6894409937888198 loss: 5.1167313  
07236478

Epoch: 5/10 Train batch:8000/31779 acc: 0.7052845528455285 loss: 5.0985987  
41926253

Epoch: 5/10 Train batch:10000/31779 acc: 0.65625 loss: 5.090086501324549

Epoch: 5/10 Train batch:12000/31779 acc: 0.6460905349794238 loss: 5.081751  
8518306315

Epoch: 5/10 Train batch:14000/31779 acc: 0.696969696969697 loss: 5.0804935  
67278609

Epoch: 5/10 Train batch:16000/31779 acc: 0.6275720164609053 loss: 5.084368  
4275168926

Epoch: 5/10 Train batch:18000/31779 acc: 0.6659836065573771 loss: 5.072955

```
843992531
Epoch: 5/10 Train batch:20000/31779 acc: 0.6511156186612576 loss: 5.063052
929705009
Epoch: 5/10 Train batch:22000/31779 acc: 0.6880165289256198 loss: 5.057854
80258055
Epoch: 5/10 Train batch:24000/31779 acc: 0.694560669456067 loss: 5.0376209
418755025
Epoch: 5/10 Train batch:26000/31779 acc: 0.6912065439672802 loss: 5.037911
129882559
Epoch: 5/10 Train batch:28000/31779 acc: 0.651356993736952 loss: 5.0399528
671987355
Epoch: 5/10 Train batch:30000/31779 acc: 0.6495901639344263 loss: 5.034081
088611856
Validation Accuracy: 0.6682242437514426. Cost time: 24.154389715194704 minute
s
=====
Epoch: 6/10 Train batch:2000/31779 acc: 0.609504132231405 loss: 4.94795658
3695486
Epoch: 6/10 Train batch:4000/31779 acc: 0.6456211812627292 loss: 4.9554243
37415025
Epoch: 6/10 Train batch:6000/31779 acc: 0.6715481171548117 loss: 4.9398640
28710872
Epoch: 6/10 Train batch:8000/31779 acc: 0.6659836065573771 loss: 4.9363483
38378593
Epoch: 6/10 Train batch:10000/31779 acc: 0.6743697478991597 loss: 4.935492
197517306
Epoch: 6/10 Train batch:12000/31779 acc: 0.6409185803757829 loss: 4.931298
549287021
Epoch: 6/10 Train batch:14000/31779 acc: 0.6340956340956341 loss: 4.934873
427497223
Epoch: 6/10 Train batch:16000/31779 acc: 0.6268041237113402 loss: 4.929511
033347808
Epoch: 6/10 Train batch:18000/31779 acc: 0.6902286902286903 loss: 4.925154
0408004075

Epoch: 6/10 Train batch:20000/31779 acc: 0.6757322175732218 loss: 4.92136750
8592084
Epoch: 6/10 Train batch:22000/31779 acc: 0.689795918367347 loss: 4.916912987
595424
Epoch: 6/10 Train batch:24000/31779 acc: 0.7276507276507277 loss: 4.91675918
2233363
Epoch: 6/10 Train batch:26000/31779 acc: 0.6735966735966736 loss: 4.89501824
0902573
Epoch: 6/10 Train batch:28000/31779 acc: 0.6604938271604939 loss: 4.89544831
4957321
Epoch: 6/10 Train batch:30000/31779 acc: 0.6997929606625258 loss: 4.89364462
3807631
Validation Accuracy: 0.6722858800063921. Cost time: 24.207170498371124 minutes
=====
Epoch: 7/10 Train batch:2000/31779 acc: 0.6473029045643154 loss: 4.823618147
28938
Epoch: 7/10 Train batch:4000/31779 acc: 0.6871035940803383 loss: 4.811582079
6228945
Epoch: 7/10 Train batch:6000/31779 acc: 0.6911764705882353 loss: 4.816671787
411906
Epoch: 7/10 Train batch:8000/31779 acc: 0.709278350515464 loss: 4.8084125587
48387
```

```
Epoch: 7/10 Train batch:10000/31779 acc: 0.676954732510288 loss: 4.810286804
2187765
Epoch: 7/10 Train batch:12000/31779 acc: 0.6514522821576764 loss: 4.80350687
2034632
Epoch: 7/10 Train batch:14000/31779 acc: 0.65439672801636 loss: 4.7988598493
39351
Epoch: 7/10 Train batch:16000/31779 acc: 0.6570841889117043 loss: 4.79110216
1863819
Epoch: 7/10 Train batch:18000/31779 acc: 0.6687116564417178 loss: 4.78965670
8016992
Epoch: 7/10 Train batch:20000/31779 acc: 0.6772823779193206 loss: 4.79083626
6707629
Epoch: 7/10 Train batch:22000/31779 acc: 0.6940928270042194 loss: 4.78916440
2987808
Epoch: 7/10 Train batch:24000/31779 acc: 0.6867219917012448 loss: 4.79491839
0456587
Epoch: 7/10 Train batch:26000/31779 acc: 0.6804123711340206 loss: 4.76477698
3647607
Epoch: 7/10 Train batch:28000/31779 acc: 0.6804979253112033 loss: 4.77703860
2274843
Epoch: 7/10 Train batch:30000/31779 acc: 0.6639175257731958 loss: 4.76454665
0694683
Validation Accuracy: 0.6772663782337727. Cost time: 24.20336433649063 minutes
=====
Epoch: 8/10 Train batch:2000/31779 acc: 0.6796714579055442 loss: 4.709960559
150204
Epoch: 8/10 Train batch:4000/31779 acc: 0.6818181818181818 loss: 4.698262932
128273
Epoch: 8/10 Train batch:6000/31779 acc: 0.675 loss: 4.705215996829793
Epoch: 8/10 Train batch:8000/31779 acc: 0.681912681912682 loss: 4.7214952277
71811
Epoch: 8/10 Train batch:10000/31779 acc: 0.6529774127310062 loss: 4.69487122
8188276
Epoch: 8/10 Train batch:12000/31779 acc: 0.6611570247933884 loss: 4.71155673
62634465
Epoch: 8/10 Train batch:14000/31779 acc: 0.6762295081967213 loss: 4.69137926
52776465
Epoch: 8/10 Train batch:16000/31779 acc: 0.7098765432098766 loss: 4.69194528
6118425
Epoch: 8/10 Train batch:18000/31779 acc: 0.7037037037037037 loss: 4.68904441
955965
Epoch: 8/10 Train batch:20000/31779 acc: 0.652542372881356 loss: 4.686819446
622394
Epoch: 8/10 Train batch:22000/31779 acc: 0.6550308008213552 loss: 4.69944124
1915338
Epoch: 8/10 Train batch:24000/31779 acc: 0.6735112936344969 loss: 4.68003176
5834428
Epoch: 8/10 Train batch:26000/31779 acc: 0.6855345911949685 loss: 4.68023906
3454792
Epoch: 8/10 Train batch:28000/31779 acc: 0.6862348178137652 loss: 4.68114329
7930248
Epoch: 8/10 Train batch:30000/31779 acc: 0.7018633540372671 loss: 4.67489990
3242476
Validation Accuracy: 0.6798705603134488. Cost time: 24.238636338710783 minutes
=====
Epoch: 9/10 Train batch:2000/31779 acc: 0.6673596673596673 loss: 4.625237669
33009
```

```
Epoch: 9/10 Train batch:4000/31779 acc: 0.6777546777546778 loss: 4.633318156
236783
Epoch: 9/10 Train batch:6000/31779 acc: 0.6894409937888198 loss: 4.631312782
526948
Epoch: 9/10 Train batch:8000/31779 acc: 0.6927835051546392 loss: 4.628271773
690358
Epoch: 9/10 Train batch:10000/31779 acc: 0.7125 loss: 4.638095643138513
Epoch: 9/10 Train batch:12000/31779 acc: 0.6900826446280992 loss: 4.62681207
1190216
Epoch: 9/10 Train batch:14000/31779 acc: 0.689795918367347 loss: 4.634401623
974554
Epoch: 9/10 Train batch:16000/31779 acc: 0.7154639175257732 loss: 4.62405539
5019241
Epoch: 9/10 Train batch:18000/31779 acc: 0.6834381551362684 loss: 4.62453045
9055677
Epoch: 9/10 Train batch:20000/31779 acc: 0.6923076923076923 loss: 4.61840778
5426825
Epoch: 9/10 Train batch:22000/31779 acc: 0.7134146341463414 loss: 4.61489241
1511391
Epoch: 9/10 Train batch:24000/31779 acc: 0.6652977412731006 loss: 4.62675437
1216521
Epoch: 9/10 Train batch:26000/31779 acc: 0.65625 loss: 4.615499716019258
Epoch: 9/10 Train batch:28000/31779 acc: 0.6956521739130435 loss: 4.62863692
2338046
Epoch: 9/10 Train batch:30000/31779 acc: 0.66875 loss: 4.606372268171981
Save one candidate model.
Validation Accuracy: 0.6809092738475014. Cost time: 24.235928801695504 minutes
=====
Epoch: 10/10 Train batch:2000/31779 acc: 0.6987704918032787 loss: 4.59963982
575573
Epoch: 10/10 Train batch:4000/31779 acc: 0.6653061224489796 loss: 4.59370297
5544147
Epoch: 10/10 Train batch:6000/31779 acc: 0.6356107660455487 loss: 4.58859146
5493664
Epoch: 10/10 Train batch:8000/31779 acc: 0.6604166666666667 loss: 4.58646045
9046066
Epoch: 10/10 Train batch:10000/31779 acc: 0.6536082474226804 loss: 4.6001926
21466704
Epoch: 10/10 Train batch:12000/31779 acc: 0.668041237113402 loss: 4.57725999
8070076
Epoch: 10/10 Train batch:14000/31779 acc: 0.6939203354297694 loss: 4.5920589
016750455
Epoch: 10/10 Train batch:16000/31779 acc: 0.6831275720164609 loss: 4.5907044
71222125
Epoch: 10/10 Train batch:18000/31779 acc: 0.6708860759493671 loss: 4.5835189
72620368
Epoch: 10/10 Train batch:20000/31779 acc: 0.6694560669456067 loss: 4.5921464
73238245
Epoch: 10/10 Train batch:22000/31779 acc: 0.6975308641975309 loss: 4.5930456
51330613
Epoch: 10/10 Train batch:24000/31779 acc: 0.6548117154811716 loss: 4.5898311
11487001
Epoch: 10/10 Train batch:26000/31779 acc: 0.6851851851851852 loss: 4.5833321
8190819
Epoch: 10/10 Train batch:28000/31779 acc: 0.7063655030800822 loss: 4.6092914
48723525
Epoch: 10/10 Train batch:30000/31779 acc: 0.717479674796748 loss: 4.58299049
```

5022386

Save one candidate model.

Validation Accuracy: 0.6819894175510035. Cost time: 24.225563311576842 minutes

In [ ]: