

```
In [1]: import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from sklearn.decomposition import PCA
import torch.utils.data
import pickle
from tqdm import tqdm
import time
import gc
import collections
```

```
In [2]: '''
set parameters
'''

base_path = '.'
train_path = base_path + '/train.npy'
test_path = base_path + '/test.npy'
train_labels_path = base_path + '/train_labels.npy'
dev_labels_path = base_path + '/dev_labels.npy'
dev_path = base_path + '/dev.npy'

pca_available = True
padding_method = 'self'
# padding_method = 'zero'

device = torch.device('cuda:0')
n_labels = 138
n_features = 40
n_epoch = 10
context_num = 18
```

```
In [3]: def load_train_data():
    t0 = time.time()
    print("Start loading training data...")
    train = np.load(train_path, allow_pickle=True)
    train_labels = np.load(train_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading training data in {0} minutes...".format(elapsed_time/60))

    return train, train_labels
```

```
In [4]: def load_validation_data():
    t0 = time.time()
    print("Start loading validation data...")
    val = np.load(dev_path, allow_pickle=True)
    val_labels = np.load(dev_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading validation data in {0} minutes...".format(elapsed_time/60))

    return val, val_labels
```

```
In [5]: def load_test_data():
    t0 = time.time()
    print("Start loading test data...")
    test = np.load(test_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading test data in {0} minutes...".format(elapsed_time/60))
    return test
```

```
In [6]: def load_and_process_data(features, labels, pca, context_num):
    """
    use the first&last frame of one utterance to pad the empty frame
    """
    t0 = time.time()
    padding_features = np.concatenate([np.concatenate(( \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=np.float32), \
                                                pca.transform(features[i]), \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=np.float32)) \
                                                for i in range(len(features)))]
    padding_features = torch.Tensor(padding_features)
    del features

    """
    corresponding label for padding frames
    """
    false_labels = np.array([-1]*context_num)
    padding_labels = np.concatenate([np.concatenate(( \
                                                false_labels, \
                                                labels[i], \
                                                false_labels)) \
                                                for i in range(len(labels)))]
    padding_labels = torch.Tensor(padding_labels)
    del labels

    gc.collect()

    return padding_features, padding_labels
```

```

In [7]: ass ContextDataset(Dataset):

    def __init__(self, context_num, features, targets):

        self.context_num = context_num
        self.features = features
        self.targets = targets

    def __len__(self):

        return len(self.targets)

    def __getitem__(self, index):

        if index-self.context_num >= 0 and index+self.context_num+1 <= len(self.targets):
            ...
            no need for padding
            ...
            X = self.features[index-self.context_num: index+self.context_num+1].res
            Y = self.targets[index].long()
        elif index-self.context_num < 0:
            ...
            padding for pre frames, actually doesnt matter since we drop this 'false'
            ...
            X = torch.cat((torch.zeros(self.context_num-index, self.features.shape[1]),
                            self.features[index-self.context_num: index].res), dim=0)
            Y = self.targets[index].long()
        else:
            ...
            padding for post frames, same as before
            ...
            X = torch.cat((self.features[index-self.context_num: index], torch.zeros(index-self.context_num, self.features.shape[1])), dim=0)
            Y = self.targets[index].long()

        return index, X, Y

```

```

In [8]: class SpeechDataset(Dataset):

    def __init__(self, speechdataset):

        self.features = [speechdataset[i][1] for i in range(len(speechdataset))]
        self.targets = [speechdataset[i][2] for i in range(len(speechdataset))]

    def __len__(self):

        return len(self.targets)

    def __getitem__(self, index):

        return index, self.features[index], self.targets[index]

```

```

In [9]: #1000,2048,1024,512,256+2,138
class SpeechNet(nn.Module):

    def __init__(self, context_num):

        super(SpeechNet, self).__init__()
        self.relu1 = nn.ReLU()
        self.relu2 = nn.ReLU()
        self.relu3 = nn.ReLU()
        self.relu4 = nn.ReLU()
        self.relu5 = nn.ReLU()
        self.relu6 = nn.ReLU()
        self.relu7 = nn.ReLU()
        self.relu8 = nn.ReLU()

        self.linear1 = nn.Linear((2*context_num+1)*pca.n_components, 2048)
        self.linear2 = nn.Linear(2048, 1024)
        self.linear3 = nn.Linear(1024, 810)
        self.linear4 = nn.Linear(810, 720)
        self.linear5 = nn.Linear(720, 512)
        self.linear6 = nn.Linear(512, 428)
        self.linear7 = nn.Linear(428, 300)
        self.linear8 = nn.Linear(300, 256)
        self.out = nn.Linear(256+2, 138)

        self.batchnorm1 = nn.BatchNorm1d(2048)
        self.batchnorm2 = nn.BatchNorm1d(1024)
        self.batchnorm3 = nn.BatchNorm1d(810)
        self.batchnorm4 = nn.BatchNorm1d(720)
        self.batchnorm5 = nn.BatchNorm1d(512)
        self.batchnorm6 = nn.BatchNorm1d(428)
        self.batchnorm7 = nn.BatchNorm1d(300)
        self.batchnorm8 = nn.BatchNorm1d(256+2)

        self.dropout1 = nn.Dropout(0.1)
        self.dropout2 = nn.Dropout(0.05)
        self.dropout3 = nn.Dropout(0.1)
        self.dropout4 = nn.Dropout(0.05)
        self.dropout5 = nn.Dropout(0.05)
        self.dropout6 = nn.Dropout(0.05)
        self.dropout7 = nn.Dropout(0.05)

    def forward(self, x):

        x = self.linear1(x)
        x = self.batchnorm1(x)
        x = self.relu1(x)
        #2048
        x = self.dropout1(x)
        x = self.linear2(x)
        x = self.batchnorm2(x)
        x = self.relu2(x)
        #1024
        x = self.dropout2(x)
        x = self.linear3(x)

```

```

x = self.batchnorm3(x)
x = self.relu3(x)
#810
x = self.dropout3(x)
x = self.linear4(x)
x = self.batchnorm4(x)
x = self.relu4(x)

#512
x = self.dropout4(x)
x = self.linear5(x)
x = self.batchnorm5(x)
x = self.relu5(x)

#512
x = self.dropout5(x)
x = self.linear6(x)
x = self.batchnorm6(x)
x = self.relu6(x)
#428
x = self.dropout6(x)
x = self.linear7(x)
x = self.batchnorm7(x)
x = self.relu7(x)
#300
x = self.dropout7(x)
x = self.linear8(x)
#300

avg_pool1 = torch.mean(x, 1, keepdims = True)
max_pool1,_ = torch.max(x, 1, keepdims = True)

conc = torch.cat((x, avg_pool1, max_pool1), 1)
conc = self.batchnorm8(conc)
output = self.out(conc)

return output

```

```

In [10]: def generate_dataset(context_num, features, labels):
    t0 = time.time()
    print("It may takes 20 minutes to generate train dataset...")
    context_dataset = ContextDataset(context_num, features, labels)
    dataset = SpeechDataset(context_dataset)
    t1 = time.time()
    print("Dataset generated. Elapsed time: {0}".format((t1-t0)/60))
    return dataset

```

```

In [11]: def weights_init(m):
    if isinstance(m, nn.Conv2d):
        xavier(m.weight.data)
        xavier(m.bias.data)

```

```
In [12]: def scale_cos(x):  
    start = 5e-3  
    end = 1e-5  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [13]: def second_scale_cos(x):  
    start = 1e-4  
    end = 1e-8  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [14]: class ParamScheduler:  
  
    def __init__(self, optimizer, scale_fn, total_steps):  
  
        self.optimizer = optimizer  
        self.scale_fn = scale_fn  
        self.total_steps = total_steps  
        self.current_iteration = 0  
  
    def batch_step(self):  
        for param_group in self.optimizer.param_groups:  
            param_group['lr'] = self.scale_fn(self.current_iteration/self.total_steps)  
  
        self.current_iteration += 1
```

```

In [15]: def train_model(train_dataloader, val_dataloader, n_epochs = 15):

    model = SpeechNet(context_num).to(device)
    model.apply(weights_init)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr = 0.001)

    '''
    set scheduler for decaying learning rate
    '''

    parameter_scheduler = ParamScheduler(optimizer, scale_cos, n_epoch*len(train_dataloader))
    candidate_model = 1
    print('Start training...')
    for i in range(n_epochs):

        t0 = time.time()

        avg_loss_1000_batch = 0
        val_correct = 0
        val_predicted = 0
        model.train()

        for index, (idx, features, labels) in enumerate(train_dataloader):

            optimizer.zero_grad()
            mask = [i for i in range(len(labels)) if labels[i] != torch.Tensor([0])]
            features = features[mask].cuda()
            labels = labels[mask].cuda()
            '''
            forward and backward
            '''

            output = model(features)
            loss = criterion(output, labels.long())
            avg_loss_1000_batch += loss.item()
            loss.backward()

            parameter_scheduler.batch_step()
            optimizer.step()

            if index % 2000 == 0 and index != 0:

                predictions = torch.max(output.data, 1)[1]
                predicted = len(features)
                correct = int(sum(predictions == labels.to(device)).cpu())
                print("Epoch: {0}/{1} Train batch:{2}/{3}  acc: {4}  loss: {5}"
                      .format(i, len(train_dataloader), len(features), len(labels), correct, loss.item()))

                avg_loss_1000_batch = 0

        for index, (idx, val_features, val_labels) in enumerate(val_dataloader):

            mask = [i for i in range(len(val_labels)) if val_labels[i] != torch.Tensor([0])]
            val_features = val_features[mask].to(device)
            val_labels = val_labels[mask].to(device)

```

```
model.eval()
outputs = model(val_features)
predictions = torch.max(outputs.data, 1)[1]
val_predicted += len(val_features)
val_correct += sum(predictions == val_labels.to(device))

epoch_acc = int(val_correct.cpu())/val_predicted
if epoch_acc >= 0.70:
    pickle.dump(model, open("candidate_model_{0}.pkl".format(candidate_model), "a"))
    print("Save one candidate model.")
    candidate_model += 1

t1 = time.time()
print("Validation Accuracy: {0}. Cost time: {1} minutes".format(int(val_correct.cpu()), t1 - t0))
print("=====")

return model
```



```

In [*]: if __name__ == '__main__':

    train_features, train_labels = load_train_data()
    val_features, val_labels = load_validation_data()
    test_features = load_test_data()

    if pca_available:
        '''
        load local pca file
        '''
        pca = pickle.load(open('pca.pkl', 'rb'))
    else:
        '''
        10 features will be enough
        '''
        pca = PCA(10).fit(np.concatenate(train_features))
        pickle.dump(pca, open('pca_{0}_features.pkl'.format(pca.n_components), 'w'))

    train_features, train_labels = load_and_process_data(train_features, train_labels)
    val_features, val_labels = load_and_process_data(val_features, val_labels, pca)

    train_context_dataset = ContextDataset(context_num, train_features, train_labels)
    val_context_dataset = ContextDataset(context_num, val_features, val_labels)

    '''
    It may takes more than 20 minutes to process since it loops over all 15 million samples
    But it could speed up later dataloader process
    And I could save this data by pickle
    The drawback is it's not flexible to feature engineering, like change context number
    '''
    '''
    train_dataset = SpeechDataset(train_context_dataset)
    val_dataset = SpeechDataset(val_context_dataset)
    '''
    '''
    train_mask = (train_dataset.targets.numpy() != -1)*1
    train_sampler = WeightedRandomSampler(weights=train_mask, num_samples=int(train_mask.sum()))
    val_mask = (val_dataset.targets.numpy() != -1)*1
    val_sampler = WeightedRandomSampler(weights=val_mask, num_samples=int(val_mask.sum()))
    '''
    '''
    train_dataloader = DataLoader(train_dataset,
                                  shuffle = True,
                                  batch_size = 512,
                                  num_workers = 0,
                                  pin_memory = True)

    val_dataloader = DataLoader(val_dataset,
                                shuffle = True,
                                batch_size = 512,
                                num_workers = 0,
                                pin_memory = True)

    '''
    '''
    Start train model.
    15 epochs by default.

```

Cost about 1 hour.
...

```
train_dataloader = DataLoader(train_context_dataset,
                              shuffle = True,
                              batch_size = 512,
                              num_workers = 0,
                              pin_memory = True)

val_dataloader = DataLoader(val_context_dataset,
                             shuffle = True,
                             batch_size = 512,
                             num_workers = 0,
                             pin_memory = True)

model = train_model(train_dataloader, val_dataloader)

pickle.dump(model, open("submission_model_68.pkl", "wb"))
#make_submission(model, test_features)
```

Start loading training data...

Done loading training data in 0.7967264811197917 minutes...

Start loading validation data...

Done loading validation data in 0.04981781244277954 minutes...

Start loading test data...

Done loading test data in 0.01579108238220215 minutes...

Start training...

Epoch: 1/15 Train batch:2000/31779 acc: 0.42592592592592593 loss: 9.775534
469401464

Epoch: 1/15 Train batch:4000/31779 acc: 0.49269311064718163 loss: 8.109487
411100417

Epoch: 1/15 Train batch:6000/31779 acc: 0.5387840670859538 loss: 7.5979077
54126936

Epoch: 1/15 Train batch:8000/31779 acc: 0.5245901639344263 loss: 7.2745072
58553058

Epoch: 1/15 Train batch:10000/31779 acc: 0.5661157024793388 loss: 7.060203
752014786

Epoch: 1/15 Train batch:12000/31779 acc: 0.5605749486652978 loss: 6.872567
464131862

Epoch: 1/15 Train batch:14000/31779 acc: 0.556935817805383 loss: 6.7424005
856737494

Epoch: 1/15 Train batch:16000/31779 acc: 0.5091649694501018 loss: 6.622533
0552551895

Epoch: 1/15 Train batch:18000/31779 acc: 0.5583333333333333 loss: 6.509914
910653606

Epoch: 1/15 Train batch:20000/31779 acc: 0.5488565488565489 loss: 6.422984
649892896

Epoch: 1/15 Train batch:22000/31779 acc: 0.5854166666666667 loss: 6.343326
085712761

Epoch: 1/15 Train batch:24000/31779 acc: 0.5422680412371134 loss: 6.289215
44062905

Epoch: 1/15 Train batch:26000/31779 acc: 0.5877551020408164 loss: 6.231776
090338826

Epoch: 1/15 Train batch:28000/31779 acc: 0.6270491803278688 loss: 6.169563
454110175

Epoch: 1/15 Train batch:30000/31779 acc: 0.5487804878048781 loss: 6.127935
012802482

Validation Accuracy: 0.6123319858663936. Cost time: 29.898462855815886 minute

S

```
=====
Epoch: 2/15 Train batch:2000/31779 acc: 0.5463917525773195 loss: 6.0177660
41215509
Epoch: 2/15 Train batch:4000/31779 acc: 0.6326530612244898 loss: 5.9596305
86439744
Epoch: 2/15 Train batch:6000/31779 acc: 0.5859213250517599 loss: 5.9484914
7181958
Epoch: 2/15 Train batch:8000/31779 acc: 0.6008316008316008 loss: 5.9084361
93363741
Epoch: 2/15 Train batch:10000/31779 acc: 0.5873684210526315 loss: 5.887612
983118743
Epoch: 2/15 Train batch:12000/31779 acc: 0.6376518218623481 loss: 5.862514
947075397
Epoch: 2/15 Train batch:14000/31779 acc: 0.6153846153846154 loss: 5.820339
681347832
Epoch: 2/15 Train batch:16000/31779 acc: 0.5941422594142259 loss: 5.798072
6168025285
Epoch: 2/15 Train batch:18000/31779 acc: 0.6137787056367432 loss: 5.779028
256423771
Epoch: 2/15 Train batch:20000/31779 acc: 0.6205450733752621 loss: 5.765704
936115071
Epoch: 2/15 Train batch:22000/31779 acc: 0.6296296296296297 loss: 5.736493
79843846
Epoch: 2/15 Train batch:24000/31779 acc: 0.6112266112266113 loss: 5.712157
631292939
Epoch: 2/15 Train batch:26000/31779 acc: 0.6131687242798354 loss: 5.691869
272151962
Epoch: 2/15 Train batch:28000/31779 acc: 0.5904365904365905 loss: 5.677539
088530466
Epoch: 2/15 Train batch:30000/31779 acc: 0.629399585921325 loss: 5.6451472
66564891
Validation Accuracy: 0.6391254683088796. Cost time: 23.292837572097778 minute
```

S

```
=====
Epoch: 3/15 Train batch:2000/31779 acc: 0.6145833333333334 loss: 5.5592634
84319672
Epoch: 3/15 Train batch:4000/31779 acc: 0.6155419222903885 loss: 5.5491820
180322975
Epoch: 3/15 Train batch:6000/31779 acc: 0.6114519427402862 loss: 5.5398837
50258014
Epoch: 3/15 Train batch:8000/31779 acc: 0.6354166666666666 loss: 5.5209490
56139216
Epoch: 3/15 Train batch:10000/31779 acc: 0.6088709677419355 loss: 5.514790
925895795
Epoch: 3/15 Train batch:12000/31779 acc: 0.6526315789473685 loss: 5.489247
277611867
Epoch: 3/15 Train batch:14000/31779 acc: 0.5950413223140496 loss: 5.482605
571858585
Epoch: 3/15 Train batch:16000/31779 acc: 0.6350515463917525 loss: 5.471510
59191674
Epoch: 3/15 Train batch:18000/31779 acc: 0.6086956521739131 loss: 5.449926
839210093
Epoch: 3/15 Train batch:20000/31779 acc: 0.5955284552845529 loss: 5.452208
753209561
Epoch: 3/15 Train batch:22000/31779 acc: 0.6464435146443515 loss: 5.437912
923749536
```

```
Epoch: 3/15 Train batch:24000/31779 acc: 0.6638830897703549 loss: 5.425580
8622110635
Epoch: 3/15 Train batch:26000/31779 acc: 0.6523517382413088 loss: 5.404229
0057986975
Epoch: 3/15 Train batch:28000/31779 acc: 0.6028513238289206 loss: 5.386188
2362980396
Epoch: 3/15 Train batch:30000/31779 acc: 0.6157112526539278 loss: 5.392457
937821746
Validation Accuracy: 0.652437869542315. Cost time: 24.991123350461326 minutes
=====
Epoch: 4/15 Train batch:2000/31779 acc: 0.6962809917355371 loss: 5.3062039
02699053
Epoch: 4/15 Train batch:4000/31779 acc: 0.6473029045643154 loss: 5.3037068
39447841
Epoch: 4/15 Train batch:6000/31779 acc: 0.6465696465696466 loss: 5.2949280
78074008
Epoch: 4/15 Train batch:8000/31779 acc: 0.6409185803757829 loss: 5.2836531
86634183
Epoch: 4/15 Train batch:10000/31779 acc: 0.6659751037344398 loss: 5.281832
566251978
Epoch: 4/15 Train batch:12000/31779 acc: 0.6452282157676349 loss: 5.256409
40410085
Epoch: 4/15 Train batch:14000/31779 acc: 0.6680497925311203 loss: 5.251310
472376645
Epoch: 4/15 Train batch:16000/31779 acc: 0.6476578411405295 loss: 5.252602
803520858
Epoch: 4/15 Train batch:18000/31779 acc: 0.6273684210526316 loss: 5.243135
431781411
Epoch: 4/15 Train batch:20000/31779 acc: 0.6751054852320675 loss: 5.246580
079197884
Epoch: 4/15 Train batch:22000/31779 acc: 0.6520833333333333 loss: 5.220654
300646856
Epoch: 4/15 Train batch:24000/31779 acc: 0.6556016597510373 loss: 5.222006
601979956
Epoch: 4/15 Train batch:26000/31779 acc: 0.6468172484599589 loss: 5.219035
235699266
Epoch: 4/15 Train batch:28000/31779 acc: 0.620253164556962 loss: 5.1856239
55447227
Epoch: 4/15 Train batch:30000/31779 acc: 0.6836734693877551 loss: 5.192454
103147611
Validation Accuracy: 0.6618336401138738. Cost time: 32.523609634240465 minute
s
=====
Epoch: 5/15 Train batch:2000/31779 acc: 0.6625258799171843 loss: 5.1132215
19626677
Epoch: 5/15 Train batch:4000/31779 acc: 0.6756198347107438 loss: 5.1009923
00081998
Epoch: 5/15 Train batch:6000/31779 acc: 0.6551020408163265 loss: 5.1157691
72552973
Epoch: 5/15 Train batch:8000/31779 acc: 0.694672131147541 loss: 5.09827569
5927441
Epoch: 5/15 Train batch:10000/31779 acc: 0.6826722338204593 loss: 5.089762
9777900875
Epoch: 5/15 Train batch:12000/31779 acc: 0.6700819672131147 loss: 5.092070
585116744
Epoch: 5/15 Train batch:14000/31779 acc: 0.6412371134020619 loss: 5.075249
331071973
```

```
Epoch: 5/15 Train batch:16000/31779 acc: 0.6638655462184874 loss: 5.076793
9439974725
Epoch: 5/15 Train batch:18000/31779 acc: 0.6371134020618556 loss: 5.072140
499367379
Epoch: 5/15 Train batch:20000/31779 acc: 0.663135593220339 loss: 5.0696268
0327706
Epoch: 5/15 Train batch:22000/31779 acc: 0.6784968684759917 loss: 5.057284
829672426
Epoch: 5/15 Train batch:24000/31779 acc: 0.7027027027027027 loss: 5.051002
2402741015
Epoch: 5/15 Train batch:26000/31779 acc: 0.6375266524520256 loss: 5.050701
953005046
Epoch: 5/15 Train batch:28000/31779 acc: 0.6659707724425887 loss: 5.041448
8753303885
Epoch: 5/15 Train batch:30000/31779 acc: 0.6155419222903885 loss: 5.019512
769533321
Validation Accuracy: 0.6674311519362686. Cost time: 29.610111657778422 minute
s
=====
Epoch: 6/15 Train batch:2000/31779 acc: 0.6673511293634496 loss: 4.9649889
0966177
Epoch: 6/15 Train batch:4000/31779 acc: 0.6199186991869918 loss: 4.9560458
56233686
Epoch: 6/15 Train batch:6000/31779 acc: 0.6503067484662577 loss: 4.9370094
67743337
Epoch: 6/15 Train batch:8000/31779 acc: 0.6776180698151951 loss: 4.9488178
29135805
Epoch: 6/15 Train batch:10000/31779 acc: 0.6434426229508197 loss: 4.938042
079564184
Epoch: 6/15 Train batch:12000/31779 acc: 0.6474226804123712 loss: 4.938233
001390472
Epoch: 6/15 Train batch:14000/31779 acc: 0.6750524109014675 loss: 4.937296
177376993
Epoch: 6/15 Train batch:16000/31779 acc: 0.6406570841889117 loss: 4.921966
823982075

Epoch: 6/15 Train batch:18000/31779 acc: 0.6873706004140787 loss: 4.91727610
3980839
Epoch: 6/15 Train batch:20000/31779 acc: 0.6811594202898551 loss: 4.91769523
2201368
Epoch: 6/15 Train batch:22000/31779 acc: 0.6680327868852459 loss: 4.90579132
8288615
Epoch: 6/15 Train batch:24000/31779 acc: 0.6687116564417178 loss: 4.89853619
8306829
Epoch: 6/15 Train batch:26000/31779 acc: 0.6761710794297352 loss: 4.89721436
8451387
Epoch: 6/15 Train batch:28000/31779 acc: 0.6673596673596673 loss: 4.89238152
3037329
Epoch: 6/15 Train batch:30000/31779 acc: 0.676954732510288 loss: 4.896091582
486406
Validation Accuracy: 0.6729739167490338. Cost time: 24.34170482158661 minutes
=====
Epoch: 7/15 Train batch:2000/31779 acc: 0.6639344262295082 loss: 4.816347099
957056
Epoch: 7/15 Train batch:4000/31779 acc: 0.6604938271604939 loss: 4.815884839
161299
Epoch: 7/15 Train batch:6000/31779 acc: 0.6728778467908902 loss: 4.811424470
```

8558545

Epoch: 7/15 Train batch:8000/31779 acc: 0.6521739130434783 loss: 4.817386562
121101

Epoch: 7/15 Train batch:10000/31779 acc: 0.6447638603696099 loss: 4.80826326
9718736

Epoch: 7/15 Train batch:12000/31779 acc: 0.691683569979716 loss: 4.808936178
451404

Epoch: 7/15 Train batch:14000/31779 acc: 0.6687370600414079 loss: 4.80394072
5745633

Epoch: 7/15 Train batch:16000/31779 acc: 0.6413934426229508 loss: 4.79524684
0920299

Epoch: 7/15 Train batch:18000/31779 acc: 0.6556701030927835 loss: 4.79816957
7028602

In []: