

```
In [1]: import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from sklearn.decomposition import PCA
import torch.utils.data
import pickle
from tqdm import tqdm
import time
import gc
import collections
```

```
In [2]: '''
set parameters
'''

base_path = '.'
train_path = base_path + '/train.npy'
test_path = base_path + '/test.npy'
train_labels_path = base_path + '/train_labels.npy'
dev_labels_path = base_path + '/dev_labels.npy'
dev_path = base_path + '/dev.npy'

pca_available = True
padding_method = 'self'
# padding_method = 'zero'

device = torch.device('cuda:0')
n_labels = 138
n_features = 40
n_epoch = 10
context_num = 18
```

```
In [3]: def load_train_data():
    t0 = time.time()
    print("Start loading training data...")
    train = np.load(train_path, allow_pickle=True)
    train_labels = np.load(train_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading training data in {0} minutes...".format(elapsed_time/60))

    return train, train_labels
```

```
In [4]: def load_validation_data():
    t0 = time.time()
    print("Start loading validation data...")
    val = np.load(dev_path, allow_pickle=True)
    val_labels = np.load(dev_labels_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading validation data in {0} minutes...".format(elapsed_time/60))

    return val, val_labels
```

```
In [5]: def load_test_data():
    t0 = time.time()
    print("Start loading test data...")
    test = np.load(test_path, allow_pickle=True)
    t1 = time.time()
    elapsed_time = t1 - t0
    print("Done loading test data in {0} minutes...".format(elapsed_time/60))
    return test
```

```
In [6]: def load_and_process_data(features, labels, pca, context_num):
    """
    use the first&last frame of one utterance to pad the empty frame
    """
    t0 = time.time()
    padding_features = np.concatenate([np.concatenate(( \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=pca.dtype), \
                                                pca.transform(features[i]), \
                                                np.ones((context_num, pca.n_comps), \
                                                        dtype=pca.dtype)) \
                                                for i in range(len(features)))]
    padding_features = torch.Tensor(padding_features)
    del features

    """
    corresponding label for padding frames
    """
    false_labels = np.array([-1]*context_num)
    padding_labels = np.concatenate([np.concatenate(( \
                                                false_labels, \
                                                labels[i], \
                                                false_labels)) \
                                                for i in range(len(labels)))]
    padding_labels = torch.Tensor(padding_labels)
    del labels

    gc.collect()

    return padding_features, padding_labels
```

```

In [7]: class ContextDataset(Dataset):

    def __init__(self, context_num, features, targets):

        self.context_num = context_num
        self.features = features
        self.targets = targets

    def __len__(self):

        return len(self.targets)

    def __getitem__(self, index):

        if index-self.context_num >= 0 and index+self.context_num+1 <= len(self.
            ...
            no need for padding
            ...

            X = self.features[index-self.context_num: index+self.context_num+1].
            Y = self.targets[index].long()
        elif index-self.context_num < 0:
            ...
            padding for pre frames, actually doesnt matter since we drop this 'f
            ...

            X = torch.cat((torch.zeros(self.context_num-index, self.features.sha
            Y = self.targets[index].long()
        else:
            ...
            padding for post frames, same as before
            ...

            X = torch.cat((self.features[index-self.context_num:], torch.zeros(i
            Y = self.targets[index].long()

        return index, X, Y

```

```

In [8]: #1000,2048,1024,512,256+2,138
class SpeechNet(nn.Module):

    def __init__(self, context_num):

        super(SpeechNet, self).__init__()
        self.relu1 = nn.ReLU()
        self.relu2 = nn.ReLU()
        self.relu3 = nn.ReLU()
        self.relu4 = nn.ReLU()
        self.relu5 = nn.ReLU()
        self.relu6 = nn.ReLU()
        self.relu7 = nn.ReLU()
        self.relu8 = nn.ReLU()

        self.linear1 = nn.Linear((2*context_num+1)*pca.n_components, 2048)
        self.linear2 = nn.Linear(2048, 1024)
        self.linear3 = nn.Linear(1024, 810)
        self.linear4 = nn.Linear(810, 720)
        self.linear5 = nn.Linear(720, 512)
        self.linear6 = nn.Linear(512, 428)
        self.linear7 = nn.Linear(428, 300)
        self.linear8 = nn.Linear(300, 256)
        self.out = nn.Linear(256+2, 138)

        self.batchnorm1 = nn.BatchNorm1d(2048)
        self.batchnorm2 = nn.BatchNorm1d(1024)
        self.batchnorm3 = nn.BatchNorm1d(810)
        self.batchnorm4 = nn.BatchNorm1d(720)
        self.batchnorm5 = nn.BatchNorm1d(512)
        self.batchnorm6 = nn.BatchNorm1d(428)
        self.batchnorm7 = nn.BatchNorm1d(300)
        self.batchnorm8 = nn.BatchNorm1d(256+2)

        self.dropout1 = nn.Dropout(0.1)
        self.dropout2 = nn.Dropout(0.05)
        self.dropout3 = nn.Dropout(0.1)
        self.dropout4 = nn.Dropout(0.05)
        self.dropout5 = nn.Dropout(0.05)
        self.dropout6 = nn.Dropout(0.05)
        self.dropout7 = nn.Dropout(0.05)

    def forward(self, x):

        x = self.linear1(x)
        x = self.batchnorm1(x)
        x = self.relu1(x)
        #2048
        x = self.dropout1(x)
        x = self.linear2(x)
        x = self.batchnorm2(x)
        x = self.relu2(x)
        #1024
        x = self.dropout2(x)
        x = self.linear3(x)

```

```

x = self.batchnorm3(x)
x = self.relu3(x)
#810
x = self.dropout3(x)
x = self.linear4(x)
x = self.batchnorm4(x)
x = self.relu4(x)

#512
x = self.dropout4(x)
x = self.linear5(x)
x = self.batchnorm5(x)
x = self.relu5(x)

#512
x = self.dropout5(x)
x = self.linear6(x)
x = self.batchnorm6(x)
x = self.relu6(x)
#428
x = self.dropout6(x)
x = self.linear7(x)
x = self.batchnorm7(x)
x = self.relu7(x)
#300
x = self.dropout7(x)
x = self.linear8(x)
#300

avg_pool1 = torch.mean(x, 1, keepdims = True)
max_pool1,_ = torch.max(x, 1, keepdims = True)

conc = torch.cat((x, avg_pool1, max_pool1), 1)
conc = self.batchnorm8(conc)
output = self.out(conc)

return output

```

```

In [9]: def generate_dataset(context_num, features, labels):
        t0 = time.time()
        print("It may takes 20 minutes to generate train dataset...")
        context_dataset = ContextDataset(context_num, features, labels)
        dataset = SpeechDataset(context_dataset)
        t1 = time.time()
        print("Dataset generated. Elapsed time: {0}".format((t1-t0)/60))
        return dataset

```

```

In [10]: def weights_init(m):
        if isinstance(m, nn.Conv2d):
            xavier(m.weight.data)
            xavier(m.bias.data)

```

```
In [11]: def scale_cos(x):  
    start = 5e-3  
    end = 1e-5  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [12]: def second_scale_cos(x):  
    start = 1e-4  
    end = 1e-8  
    return start + (1 + np.cos(np.pi * (1 - x))) * (end - start) / 2
```

```
In [13]: class ParamScheduler:  
  
    def __init__(self, optimizer, scale_fn, total_steps):  
  
        self.optimizer = optimizer  
        self.scale_fn = scale_fn  
        self.total_steps = total_steps  
        self.current_iteration = 0  
  
    def batch_step(self):  
        for param_group in self.optimizer.param_groups:  
            param_group['lr'] = self.scale_fn(self.current_iteration/self.total_steps)  
  
        self.current_iteration += 1
```

```

In [14]: def train_model(train_dataloader, val_dataloader, n_epochs = 10):

    model = SpeechNet(context_num).to(device)
    model.apply(weights_init)
    criterion = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr = 0.001)

    '''
    set scheduler for decaying learning rate
    '''

    parameter_scheduler = ParamScheduler(optimizer, scale_cos, n_epoch*len(train_dataloader))
    candidate_model = 1
    print('Start training...')
    for i in range(n_epochs):

        t0 = time.time()

        avg_loss_1000_batch = 0
        val_correct = 0
        val_predicted = 0
        model.train()

        for index, (idx, features, labels) in enumerate(train_dataloader):

            optimizer.zero_grad()
            mask = [i for i in range(len(labels)) if labels[i] != torch.Tensor([0])]
            features = features[mask].cuda()
            labels = labels[mask].cuda()
            '''
            forward and backward
            '''

            output = model(features)
            loss = criterion(output, labels.long())
            avg_loss_1000_batch += loss.item()
            loss.backward()

            parameter_scheduler.batch_step()
            optimizer.step()

            if index % 2000 == 0 and index != 0:

                predictions = torch.max(output.data, 1)[1]
                predicted = len(features)
                correct = int(sum(predictions == labels.to(device)).cpu())
                print("Epoch: {0}/{1} Train batch:{2}/{3}  acc: {4}  loss: {5}"
                      .format(i, len(train_dataloader), len(features), len(labels), correct, loss.item()))

                avg_loss_1000_batch = 0

        for index, (idx, val_features, val_labels) in enumerate(val_dataloader):

            mask = [i for i in range(len(val_labels)) if val_labels[i] != torch.Tensor([0])]
            val_features = val_features[mask].to(device)
            val_labels = val_labels[mask].to(device)

```

```
model.eval()
outputs = model(val_features)
predictions = torch.max(outputs.data, 1)[1]
val_predicted += len(val_features)
val_correct += sum(predictions == val_labels.to(device))

epoch_acc = int(val_correct.cpu())/val_predicted
if epoch_acc >= 0.68:
    pickle.dump(model, open("candidate_model_{0}.pkl".format(candidate_model), "a"))
    print("Save one candidate model.")
    candidate_model += 1

t1 = time.time()
print("Validation Accuracy: {0}. Cost time: {1} minutes".format(int(val_correct.cpu()), t1 - t0))
print("=====")

return model
```



```

In [16]: if __name__ == '__main__':

    train_features, train_labels = load_train_data()
    val_features, val_labels = load_validation_data()
    test_features = load_test_data()

    if pca_available:
        '''
        load local pca file
        '''
        pca = pickle.load(open('pca.pkl', 'rb'))
    else:
        '''
        10 features will be enough
        '''
        pca = PCA(10).fit(np.concatenate(train_features))
        pickle.dump(pca, open('pca_{0}_features.pkl'.format(pca.n_components), 'w'))

    train_features, train_labels = load_and_process_data(train_features, train_labels)
    val_features, val_labels = load_and_process_data(val_features, val_labels, pca)

    train_context_dataset = ContextDataset(context_num, train_features, train_labels)
    val_context_dataset = ContextDataset(context_num, val_features, val_labels)

    '''
    It may takes more than 20 minutes to process since it loops over all 15 million samples
    But it could speed up later dataloader process
    And I could save this data by pickle
    The drawback is it's not flexible to feature engineering, like change context number
    '''
    '''
    train_dataset = SpeechDataset(train_context_dataset)
    val_dataset = SpeechDataset(val_context_dataset)
    '''
    '''
    train_mask = (train_dataset.targets.numpy() != -1)*1
    train_sampler = WeightedRandomSampler(weights=train_mask, num_samples=int(train_mask.sum()))
    val_mask = (val_dataset.targets.numpy() != -1)*1
    val_sampler = WeightedRandomSampler(weights=val_mask, num_samples=int(val_mask.sum()))
    '''
    '''
    train_dataloader = DataLoader(train_dataset,
                                  shuffle = True,
                                  batch_size = 512,
                                  num_workers = 0,
                                  pin_memory = True)

    val_dataloader = DataLoader(val_dataset,
                                shuffle = True,
                                batch_size = 512,
                                num_workers = 0,
                                pin_memory = True)

    '''
    '''
    Start train model.

```

```
15 epochs by default.
Cost about 1 hour.
'''
```

```
train_dataloader = DataLoader(train_context_dataset,
                              shuffle = True,
                              batch_size = 512,
                              num_workers = 0,
                              pin_memory = True)

val_dataloader = DataLoader(val_context_dataset,
                             shuffle = True,
                             batch_size = 512,
                             num_workers = 0,
                             pin_memory = True)

model = train_model(train_dataloader, val_dataloader)

pickle.dump(model, open("submission_model_68.pkl", "wb"))
#make_submission(model, test_features)
```

Start loading training data...

Done loading training data in 0.7987695336341858 minutes...

Start loading validation data...

Done loading validation data in 0.0456453800201416 minutes...

Start loading test data...

Done loading test data in 0.014128851890563964 minutes...

Start training...

Epoch: 1/10 Train batch:2000/31779 acc: 0.45121951219512196 loss: 9.797406
8808835

Epoch: 1/10 Train batch:4000/31779 acc: 0.4979253112033195 loss: 8.1310059
98002365

Epoch: 1/10 Train batch:6000/31779 acc: 0.5040816326530613 loss: 7.6088751
33330002

Epoch: 1/10 Train batch:8000/31779 acc: 0.5708418891170431 loss: 7.2945281
34632856

Epoch: 1/10 Train batch:10000/31779 acc: 0.5343035343035343 loss: 7.054817
696334794

Epoch: 1/10 Train batch:12000/31779 acc: 0.5429769392033543 loss: 6.877835
999010131

Epoch: 1/10 Train batch:14000/31779 acc: 0.5979381443298969 loss: 6.748973
261332139

Epoch: 1/10 Train batch:16000/31779 acc: 0.5637860082304527 loss: 6.626700
401538983

Epoch: 1/10 Train batch:18000/31779 acc: 0.56875 loss: 6.515335530042648

Epoch: 1/10 Train batch:20000/31779 acc: 0.5379876796714579 loss: 6.443903
14957127

Epoch: 1/10 Train batch:22000/31779 acc: 0.5938775510204082 loss: 6.355136
304395273

Epoch: 1/10 Train batch:24000/31779 acc: 0.5967078189300411 loss: 6.303867
116337642

Epoch: 1/10 Train batch:26000/31779 acc: 0.5983263598326359 loss: 6.232847
114559263

Epoch: 1/10 Train batch:28000/31779 acc: 0.5479166666666667 loss: 6.187457
967782393

Epoch: 1/10 Train batch:30000/31779 acc: 0.5860655737704918 loss: 6.145368
884084746

Validation Accuracy: 0.6099275563894199. Cost time: 24.387852080663045 minute

```
S
=====
Epoch: 2/10 Train batch:2000/31779    acc: 0.6356107660455487    loss: 6.0118769
84499395
Epoch: 2/10 Train batch:4000/31779    acc: 0.5904365904365905    loss: 5.9843146
83126286
Epoch: 2/10 Train batch:6000/31779    acc: 0.6290983606557377    loss: 5.9456812
27371097
Epoch: 2/10 Train batch:8000/31779    acc: 0.5785123966942148    loss: 5.9185198
9692077
Epoch: 2/10 Train batch:10000/31779   acc: 0.6042105263157894    loss: 5.886191
22421369
Epoch: 2/10 Train batch:12000/31779   acc: 0.6125    loss: 5.852636412251741
Epoch: 2/10 Train batch:14000/31779   acc: 0.5871369294605809    loss: 5.832868
226803839
Epoch: 2/10 Train batch:16000/31779   acc: 0.5857142857142857    loss: 5.808117
415057495
Epoch: 2/10 Train batch:18000/31779   acc: 0.5995893223819302    loss: 5.789125
464623794
Epoch: 2/10 Train batch:20000/31779   acc: 0.59375    loss: 5.754060934996232
Epoch: 2/10 Train batch:22000/31779   acc: 0.6470588235294118    loss: 5.732057
94300884
Epoch: 2/10 Train batch:24000/31779   acc: 0.6504065040650406    loss: 5.722947
325091809
Epoch: 2/10 Train batch:26000/31779   acc: 0.6164948453608248    loss: 5.697432
388318703
Epoch: 2/10 Train batch:28000/31779   acc: 0.6244813278008299    loss: 5.680661
244317889
Epoch: 2/10 Train batch:30000/31779   acc: 0.6149068322981367    loss: 5.653364
775236696
Validation Accuracy: 0.6388902041323635. Cost time: 24.138027866681416 minute
s
=====
Epoch: 3/10 Train batch:2000/31779    acc: 0.6597938144329897    loss: 5.5778227
88618505
Epoch: 3/10 Train batch:4000/31779    acc: 0.6701244813278008    loss: 5.5390974
03176129
Epoch: 3/10 Train batch:6000/31779    acc: 0.6453608247422681    loss: 5.5392883
79950449
Epoch: 3/10 Train batch:8000/31779    acc: 0.6278118609406953    loss: 5.5180723
93490002
Epoch: 3/10 Train batch:10000/31779   acc: 0.6440329218106996    loss: 5.506670
6200595945
Epoch: 3/10 Train batch:12000/31779   acc: 0.6053169734151329    loss: 5.500768
842641264
Epoch: 3/10 Train batch:14000/31779   acc: 0.6529774127310062    loss: 5.487650
30503273
Epoch: 3/10 Train batch:16000/31779   acc: 0.6459627329192547    loss: 5.484318
618895486
Epoch: 3/10 Train batch:18000/31779   acc: 0.6523517382413088    loss: 5.469916
230067611
Epoch: 3/10 Train batch:20000/31779   acc: 0.6515463917525773    loss: 5.443030
263995752
Epoch: 3/10 Train batch:22000/31779   acc: 0.6040816326530613    loss: 5.427898
246562108
Epoch: 3/10 Train batch:24000/31779   acc: 0.6561181434599156    loss: 5.423127
093818039
```

```
Epoch: 3/10 Train batch:26000/31779    acc: 0.639344262295082  loss: 5.4192846
20322287
Epoch: 3/10 Train batch:28000/31779    acc: 0.6378600823045267  loss: 5.411432
175664231
Epoch: 3/10 Train batch:30000/31779    acc: 0.663135593220339  loss: 5.3958395
22359893
Validation Accuracy: 0.6519155534774709. Cost time: 24.139771298567453 minute
s
=====
Epoch: 4/10 Train batch:2000/31779    acc: 0.6257668711656442  loss: 5.3162106
93912581
Epoch: 4/10 Train batch:4000/31779    acc: 0.6548856548856549  loss: 5.2943583
80604535
Epoch: 4/10 Train batch:6000/31779    acc: 0.5907172995780591  loss: 5.2959450
49962029
Epoch: 4/10 Train batch:8000/31779    acc: 0.6321353065539113  loss: 5.2873974
24690425
Epoch: 4/10 Train batch:10000/31779   acc: 0.5917525773195876  loss: 5.273880
596272647
Epoch: 4/10 Train batch:12000/31779   acc: 0.6605691056910569  loss: 5.276953
265769407
Epoch: 4/10 Train batch:14000/31779   acc: 0.6401673640167364  loss: 5.265228
853095323
Epoch: 4/10 Train batch:16000/31779   acc: 0.6384297520661157  loss: 5.254212
558735162
Epoch: 4/10 Train batch:18000/31779   acc: 0.6346555323590815  loss: 5.241534
168599173
Epoch: 4/10 Train batch:20000/31779   acc: 0.675564681724846  loss: 5.2273399
3944712
Epoch: 4/10 Train batch:22000/31779   acc: 0.6567796610169492  loss: 5.215027
284575626
Epoch: 4/10 Train batch:24000/31779   acc: 0.6880165289256198  loss: 5.212054
678238928
Epoch: 4/10 Train batch:26000/31779   acc: 0.656964656964657  loss: 5.2214122
3680228
Epoch: 4/10 Train batch:28000/31779   acc: 0.6869918699186992  loss: 5.199652
779381722
Epoch: 4/10 Train batch:30000/31779   acc: 0.6632860040567952  loss: 5.195865
701651201
Validation Accuracy: 0.6607150255387402. Cost time: 24.175887401898702 minute
s
=====
Epoch: 5/10 Train batch:2000/31779    acc: 0.6170212765957447  loss: 5.0965385
02257317
Epoch: 5/10 Train batch:4000/31779    acc: 0.654320987654321  loss: 5.11735978
4897417
Epoch: 5/10 Train batch:6000/31779    acc: 0.6894409937888198  loss: 5.1167313
07236478
Epoch: 5/10 Train batch:8000/31779    acc: 0.7052845528455285  loss: 5.0985987
41926253
Epoch: 5/10 Train batch:10000/31779   acc: 0.65625  loss: 5.090086501324549
Epoch: 5/10 Train batch:12000/31779   acc: 0.6460905349794238  loss: 5.081751
8518306315
Epoch: 5/10 Train batch:14000/31779   acc: 0.696969696969697  loss: 5.0804935
67278609
Epoch: 5/10 Train batch:16000/31779   acc: 0.6275720164609053  loss: 5.084368
4275168926
```

Epoch: 5/10 Train batch:18000/31779 acc: 0.6659836065573771 loss: 5.072955
843992531
Epoch: 5/10 Train batch:20000/31779 acc: 0.6511156186612576 loss: 5.063052
929705009
Epoch: 5/10 Train batch:22000/31779 acc: 0.6880165289256198 loss: 5.057854
80258055
Epoch: 5/10 Train batch:24000/31779 acc: 0.694560669456067 loss: 5.0376209
418755025
Epoch: 5/10 Train batch:26000/31779 acc: 0.6912065439672802 loss: 5.037911
129882559
Epoch: 5/10 Train batch:28000/31779 acc: 0.651356993736952 loss: 5.0399528
671987355
Epoch: 5/10 Train batch:30000/31779 acc: 0.6495901639344263 loss: 5.034081
088611856
Validation Accuracy: 0.6682242437514426. Cost time: 24.154389715194704 minute
s
=====

Epoch: 6/10 Train batch:2000/31779 acc: 0.609504132231405 loss: 4.94795658
3695486
Epoch: 6/10 Train batch:4000/31779 acc: 0.6456211812627292 loss: 4.9554243
37415025
Epoch: 6/10 Train batch:6000/31779 acc: 0.6715481171548117 loss: 4.9398640
28710872
Epoch: 6/10 Train batch:8000/31779 acc: 0.6659836065573771 loss: 4.9363483
38378593
Epoch: 6/10 Train batch:10000/31779 acc: 0.6743697478991597 loss: 4.935492
197517306
Epoch: 6/10 Train batch:12000/31779 acc: 0.6409185803757829 loss: 4.931298
549287021
Epoch: 6/10 Train batch:14000/31779 acc: 0.6340956340956341 loss: 4.934873
427497223
Epoch: 6/10 Train batch:16000/31779 acc: 0.6268041237113402 loss: 4.929511
033347808
Epoch: 6/10 Train batch:18000/31779 acc: 0.6902286902286903 loss: 4.925154
0408004075
Epoch: 6/10 Train batch:20000/31779 acc: 0.6757322175732218 loss: 4.921367
508592084
Epoch: 6/10 Train batch:22000/31779 acc: 0.689795918367347 loss: 4.9169129
87595424
Epoch: 6/10 Train batch:24000/31779 acc: 0.7276507276507277 loss: 4.916759
182233363
Epoch: 6/10 Train batch:26000/31779 acc: 0.6735966735966736 loss: 4.895018
240902573
Epoch: 6/10 Train batch:28000/31779 acc: 0.6604938271604939 loss: 4.895448
314957321
Epoch: 6/10 Train batch:30000/31779 acc: 0.6997929606625258 loss: 4.893644
623807631
Validation Accuracy: 0.6722858800063921. Cost time: 24.207170498371124 minute
s
=====

Epoch: 7/10 Train batch:2000/31779 acc: 0.6473029045643154 loss: 4.8236181
4728938
Epoch: 7/10 Train batch:4000/31779 acc: 0.6871035940803383 loss: 4.8115820
796228945
Epoch: 7/10 Train batch:6000/31779 acc: 0.6911764705882353 loss: 4.8166717
87411906
Epoch: 7/10 Train batch:8000/31779 acc: 0.709278350515464 loss: 4.80841255

```
8748387
Epoch: 7/10 Train batch:10000/31779 acc: 0.676954732510288 loss: 4.8102868
042187765
Epoch: 7/10 Train batch:12000/31779 acc: 0.6514522821576764 loss: 4.803506
872034632
Epoch: 7/10 Train batch:14000/31779 acc: 0.65439672801636 loss: 4.79885984
9339351
Epoch: 7/10 Train batch:16000/31779 acc: 0.6570841889117043 loss: 4.791102
161863819
Epoch: 7/10 Train batch:18000/31779 acc: 0.6687116564417178 loss: 4.789656
708016992
Epoch: 7/10 Train batch:20000/31779 acc: 0.6772823779193206 loss: 4.790836
266707629
Epoch: 7/10 Train batch:22000/31779 acc: 0.6940928270042194 loss: 4.789164
402987808
Epoch: 7/10 Train batch:24000/31779 acc: 0.6867219917012448 loss: 4.794918
390456587
Epoch: 7/10 Train batch:26000/31779 acc: 0.6804123711340206 loss: 4.764776
983647607
Epoch: 7/10 Train batch:28000/31779 acc: 0.6804979253112033 loss: 4.777038
602274843
Epoch: 7/10 Train batch:30000/31779 acc: 0.6639175257731958 loss: 4.764546
650694683
Validation Accuracy: 0.6772663782337727. Cost time: 24.20336433649063 minutes
=====
Epoch: 8/10 Train batch:2000/31779 acc: 0.6796714579055442 loss: 4.7099605
59150204
Epoch: 8/10 Train batch:4000/31779 acc: 0.6818181818181818 loss: 4.6982629
32128273
Epoch: 8/10 Train batch:6000/31779 acc: 0.675 loss: 4.705215996829793
Epoch: 8/10 Train batch:8000/31779 acc: 0.681912681912682 loss: 4.72149522
7771811
Epoch: 8/10 Train batch:10000/31779 acc: 0.6529774127310062 loss: 4.694871
228188276
Epoch: 8/10 Train batch:12000/31779 acc: 0.6611570247933884 loss: 4.711556
7362634465
Epoch: 8/10 Train batch:14000/31779 acc: 0.6762295081967213 loss: 4.691379
2652776465
Epoch: 8/10 Train batch:16000/31779 acc: 0.7098765432098766 loss: 4.691945
286118425
Epoch: 8/10 Train batch:18000/31779 acc: 0.7037037037037037 loss: 4.689044
41955965
Epoch: 8/10 Train batch:20000/31779 acc: 0.652542372881356 loss: 4.6868194
46622394
Epoch: 8/10 Train batch:22000/31779 acc: 0.6550308008213552 loss: 4.699441
241915338
Epoch: 8/10 Train batch:24000/31779 acc: 0.6735112936344969 loss: 4.680031
765834428
Epoch: 8/10 Train batch:26000/31779 acc: 0.6855345911949685 loss: 4.680239
063454792
Epoch: 8/10 Train batch:28000/31779 acc: 0.6862348178137652 loss: 4.681143
297930248
Epoch: 8/10 Train batch:30000/31779 acc: 0.7018633540372671 loss: 4.674899
903242476
Validation Accuracy: 0.6798705603134488. Cost time: 24.238636338710783 minute
s
=====
```

```
Epoch: 9/10 Train batch:2000/31779 acc: 0.6673596673596673 loss: 4.6252376
6933009
Epoch: 9/10 Train batch:4000/31779 acc: 0.6777546777546778 loss: 4.6333181
56236783
Epoch: 9/10 Train batch:6000/31779 acc: 0.6894409937888198 loss: 4.6313127
82526948
Epoch: 9/10 Train batch:8000/31779 acc: 0.6927835051546392 loss: 4.6282717
73690358
Epoch: 9/10 Train batch:10000/31779 acc: 0.7125 loss: 4.638095643138513
Epoch: 9/10 Train batch:12000/31779 acc: 0.6900826446280992 loss: 4.626812
071190216
Epoch: 9/10 Train batch:14000/31779 acc: 0.689795918367347 loss: 4.6344016
23974554
Epoch: 9/10 Train batch:16000/31779 acc: 0.7154639175257732 loss: 4.624055
395019241
Epoch: 9/10 Train batch:18000/31779 acc: 0.6834381551362684 loss: 4.624530
459055677
Epoch: 9/10 Train batch:20000/31779 acc: 0.6923076923076923 loss: 4.618407
785426825
Epoch: 9/10 Train batch:22000/31779 acc: 0.7134146341463414 loss: 4.614892
411511391
Epoch: 9/10 Train batch:24000/31779 acc: 0.6652977412731006 loss: 4.626754
371216521
Epoch: 9/10 Train batch:26000/31779 acc: 0.65625 loss: 4.615499716019258
Epoch: 9/10 Train batch:28000/31779 acc: 0.6956521739130435 loss: 4.628636
922338046
Epoch: 9/10 Train batch:30000/31779 acc: 0.66875 loss: 4.606372268171981
Save one candidate model.
Validation Accuracy: 0.6809092738475014. Cost time: 24.235928801695504 minute
s
=====
Epoch: 10/10 Train batch:2000/31779 acc: 0.6987704918032787 loss: 4.599639
82575573
Epoch: 10/10 Train batch:4000/31779 acc: 0.6653061224489796 loss: 4.593702
975544147
Epoch: 10/10 Train batch:6000/31779 acc: 0.6356107660455487 loss: 4.588591
465493664
Epoch: 10/10 Train batch:8000/31779 acc: 0.6604166666666667 loss: 4.586460
459046066
Epoch: 10/10 Train batch:10000/31779 acc: 0.6536082474226804 loss: 4.60019
2621466704
Epoch: 10/10 Train batch:12000/31779 acc: 0.668041237113402 loss: 4.577259
998070076
Epoch: 10/10 Train batch:14000/31779 acc: 0.6939203354297694 loss: 4.59205
89016750455
Epoch: 10/10 Train batch:16000/31779 acc: 0.6831275720164609 loss: 4.59070
4471222125
Epoch: 10/10 Train batch:18000/31779 acc: 0.6708860759493671 loss: 4.58351
8972620368
Epoch: 10/10 Train batch:20000/31779 acc: 0.6694560669456067 loss: 4.59214
6473238245
Epoch: 10/10 Train batch:22000/31779 acc: 0.6975308641975309 loss: 4.59304
5651330613
Epoch: 10/10 Train batch:24000/31779 acc: 0.6548117154811716 loss: 4.58983
1111487001
Epoch: 10/10 Train batch:26000/31779 acc: 0.6851851851851852 loss: 4.58333
218190819
```

```
Epoch: 10/10 Train batch:28000/31779    acc: 0.7063655030800822  loss: 4.60929
1448723525
Epoch: 10/10 Train batch:30000/31779    acc: 0.717479674796748  loss: 4.582990
495022386
Save one candidate model.
Validation Accuracy: 0.6819894175510035. Cost time: 24.225563311576842 minute
s
=====
```

```
In [21]: def make_submission(test_features, model):

    test_false_labels = np.array([np.zeros((test_features[i].shape[0], 1)).reshape(
test_features, test_labels = load_and_process_data(test_features, test_false_labels)
'''
generate false labels(just used to skip padding frame)
'''

    test_context_dataset = ContextDataset(context_num, test_features, test_labels)
    test_dataloader = DataLoader(test_context_dataset,
                                shuffle = False,
                                batch_size = 512,
                                num_workers = 0,
                                pin_memory = True,
                                drop_last = False)

    with torch.no_grad():
        model.eval().cuda()
        sub_predictions = []
        for idx, (index, features, labels) in enumerate(test_dataloader):
            mask = [i for i in range(len(labels)) if labels[i] != torch.Tensor([
            features = features[mask].cuda()
            outputs = model(features.cuda())
            predictions = torch.max(outputs.data, 1)[1]
            sub_predictions.append(predictions)
        sub_predictions = torch.cat(sub_predictions).flatten().cpu().numpy()

    submission = pd.DataFrame(sub_predictions, columns = ['label'])
    submission.to_csv("Submission.csv", index_label = 'id')
    print("Submission.csv done.")
```



```
In [22]: if __name__ == '__main__':  
    test_features = load_test_data()  
    if pca_available:  
        '''  
        load local pca file  
        '''  
    pca = pickle.load(open('pca.pkl', 'rb'))  
    else:  
        '''  
        10 features will be enough  
        '''  
    pca = PCA(10).fit(np.concatenate(train_features))  
    pickle.dump(pca, open('pca_{0}_features.pkl'.format(pca.n_components), 'w'))  
    model = pickle.load(open("submission_model_68.pkl", "rb"))  
    make_submission(test_features, model)
```

Start loading test data...

Done loading test data in 0.005717750390370687 minutes...

Submission.csv done.

In []: